

Discussion Section 3 (CS145)

2015-10-16

Week 03

Outline

- Homework 2 will be posted on Monday (Oct 19)
 - Due Oct 26 at noon (beginning of the class)
- Review:
 - Clustering Algorithms
 - Hierarchical Clustering
 - B+Tree & BIRCH
 - DBSCAN

Why clustering

- Discover some hidden interesting patterns.
 - Example:
 - (1) abbreviation: UCLA, University of California Los Angeles
 - (2) cognates (cross-language): Vienna theater, Vienna theatre
 - (3) similar expression: Florida fine cars, Florida fine auto
 - (4) Arabic numerals vs English words: 24 hours, twenty-four hours

What is a good clustering

- (1) Keep similar objects together and dissimilar objects apart.
- (2) In other words, high ***intra-class*** similarity and low ***inter-class*** similarity.


Hierarchical Clustering

- Goal: Group data objects into a tree of clusters
- Approaches:
 - Agglomerative
 - A “bottom up” approach
 - Each object starts as its own cluster
 - A pair of clusters are merged at each iteration until all objects form a single big cluster
 - Divisive
 - A “top down” approach
 - All objects start as one big cluster
 - Clusters are split at each iteration until each cluster contains only one object

Hierarchical Clustering

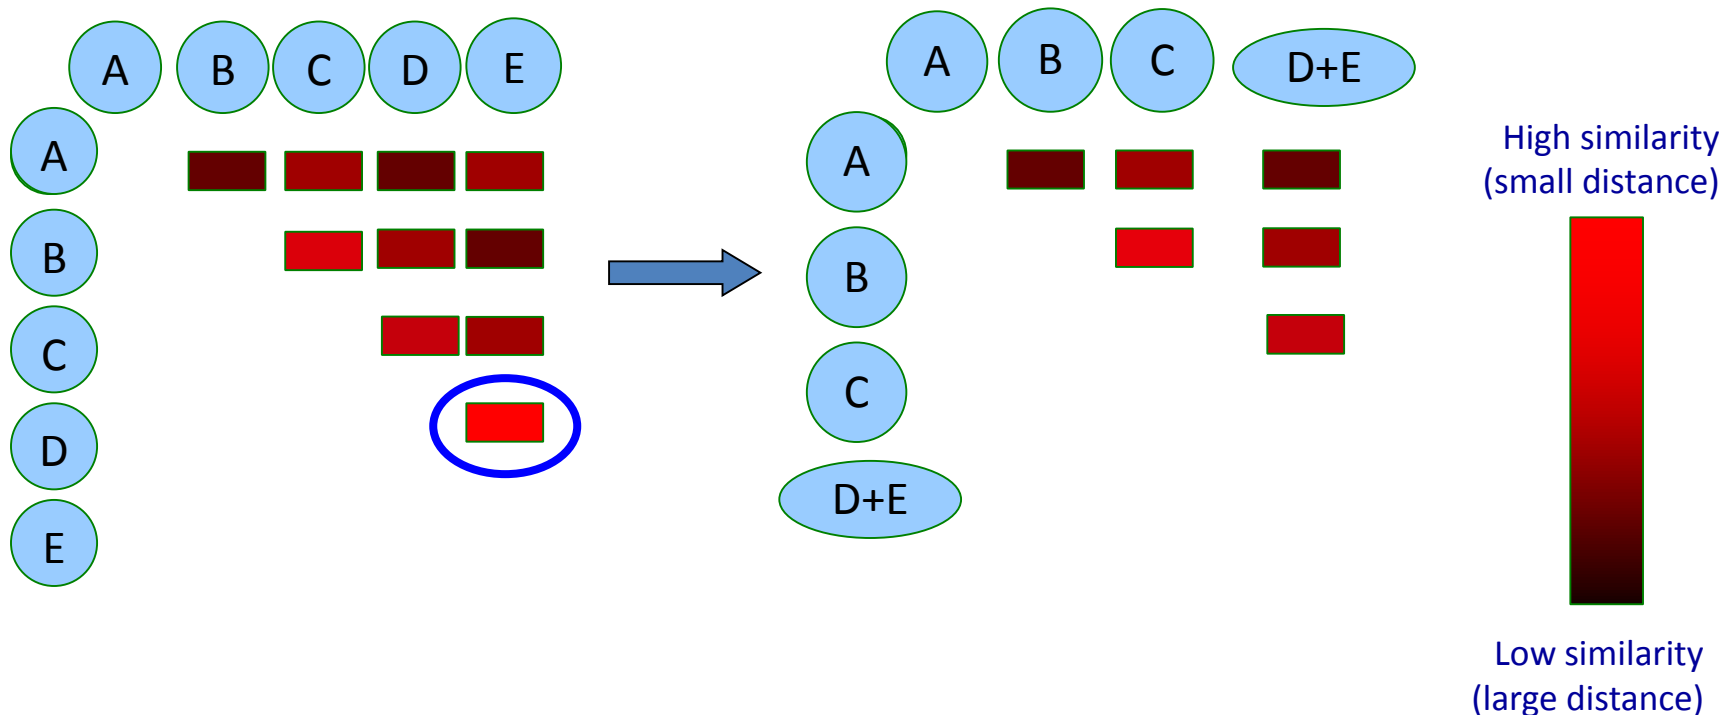
- Agglomerative Approach
 - Step 1: Calculate the distance matrix
 - Step 2: Join two members with the closest distance and recalculate the distance matrix
 - Step 3: repeat step 2 until all members form a single cluster

| | Feature X | Feature Y | Feature Z |
|---|-----------|-----------|-----------|
| A | X_A | Y_A | Z_A |
| B | X_B | Y_B | Z_B |
| C | X_C | Y_C | Z_C |
| D | X_D | Y_D | Z_D |
| E | X_E | Y_E | Z_E |



Hierarchical Clustering


- Agglomerative Approach
 - Step 1: Calculate the distance matrix
 - Step 2: Join two members with the closest distance and recalculate the distance matrix
 - Step 3: repeat step 2 until all members form a single cluster




Hierarchical Clustering

- Agglomerative Approach
 - Step 1: Calculate the distance matrix
 - Step 2: Join two members with the closest distance and recalculate the distance matrix
 - Step 3: repeat step 2 until all members form a single cluster


Distance Measure Between Clusters




1. Minimum Distance (Single Linkage)



2. Maximum Distance (Complete Linkage)



3. Mean Distance

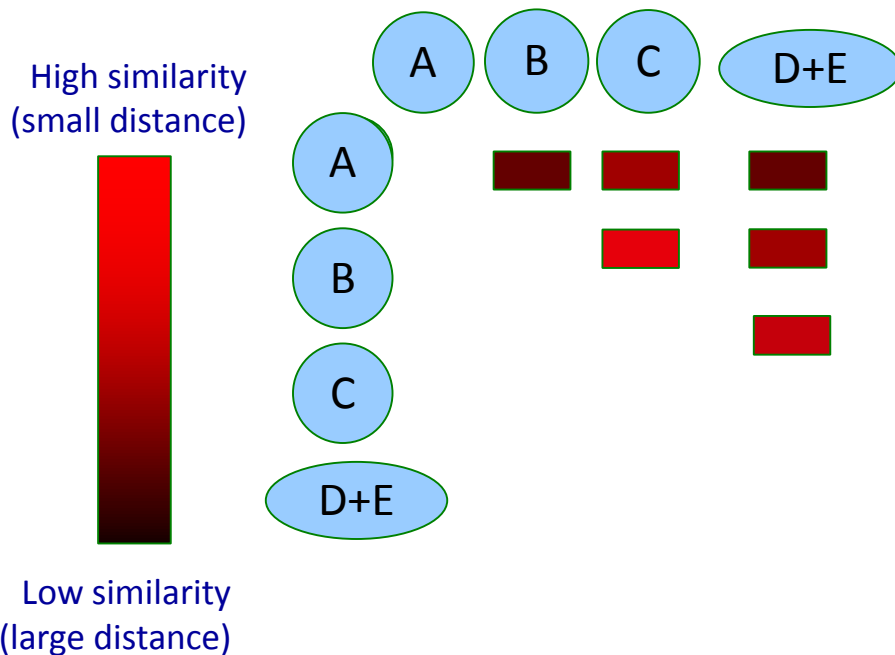


4. Average Distance


(average of all pairs)

Hierarchical Clustering

- Agglomerative Approach
 - Step 1: Calculate the distance matrix
 - Step 2: Join two members with the closest distance and recalculate the distance matrix
 - Step 3: repeat step 2 until all members form a single cluster



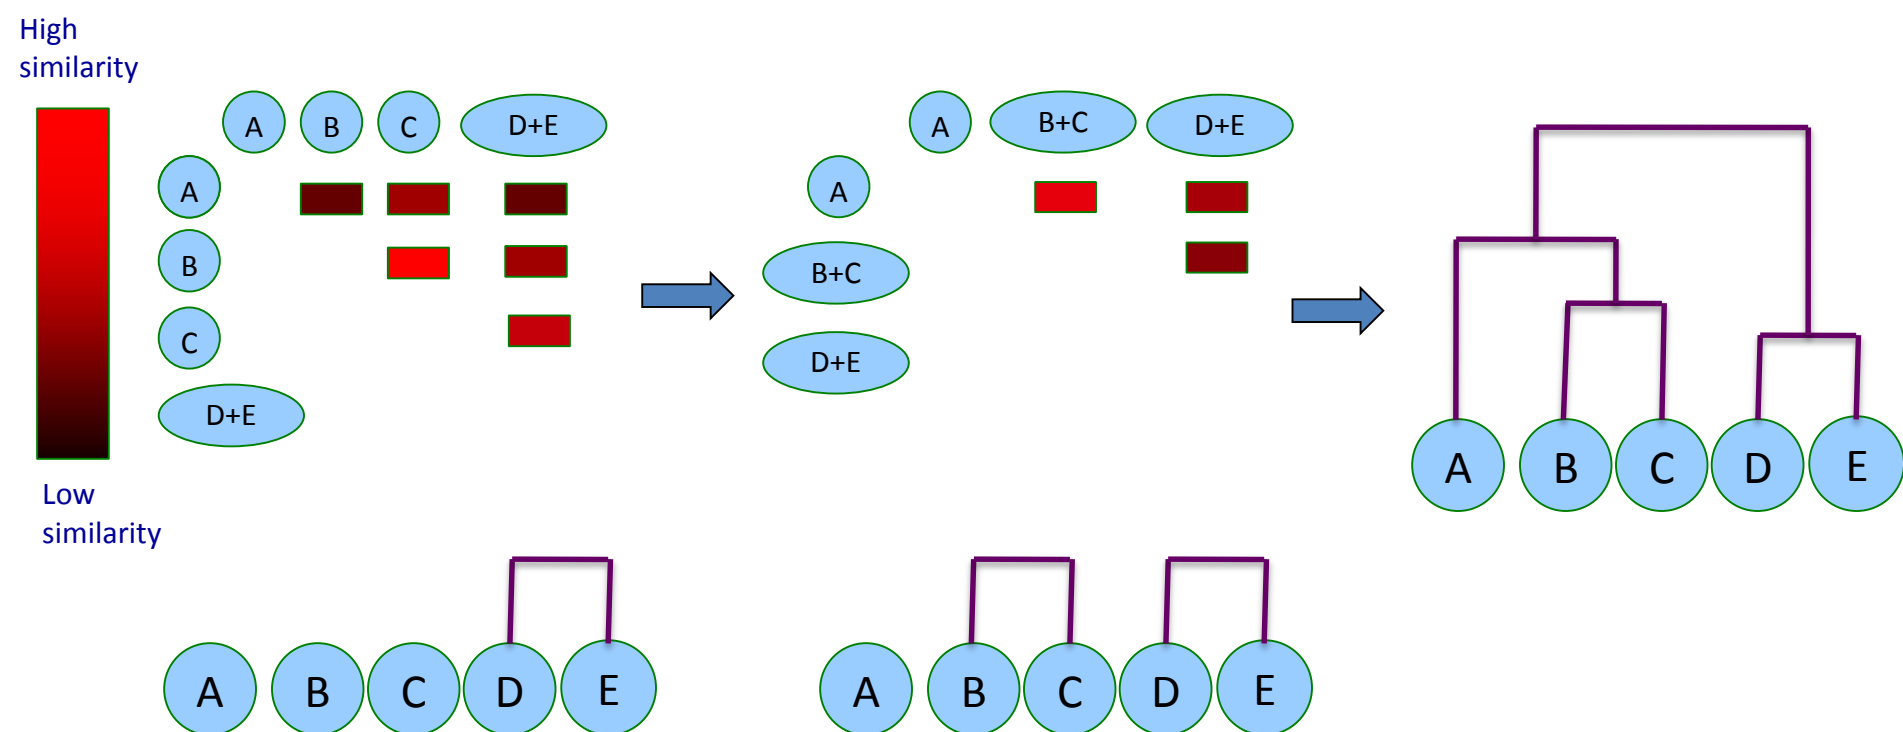
Dendrogram



Hierarchical Clustering

- Agglomerative Approach

- Step 1: Calculate the distance matrix
- Step 2: Join two members with the closest distance and recalculate the distance matrix
- Step 3: repeat step 2 until all members form a single cluster



Hierarchical Clustering

- Challenges
 - Hard to choose merge/split points
 - Never undo merging/splitting
 - Do not scale well
 - Data may not fit in memory

BIRCH

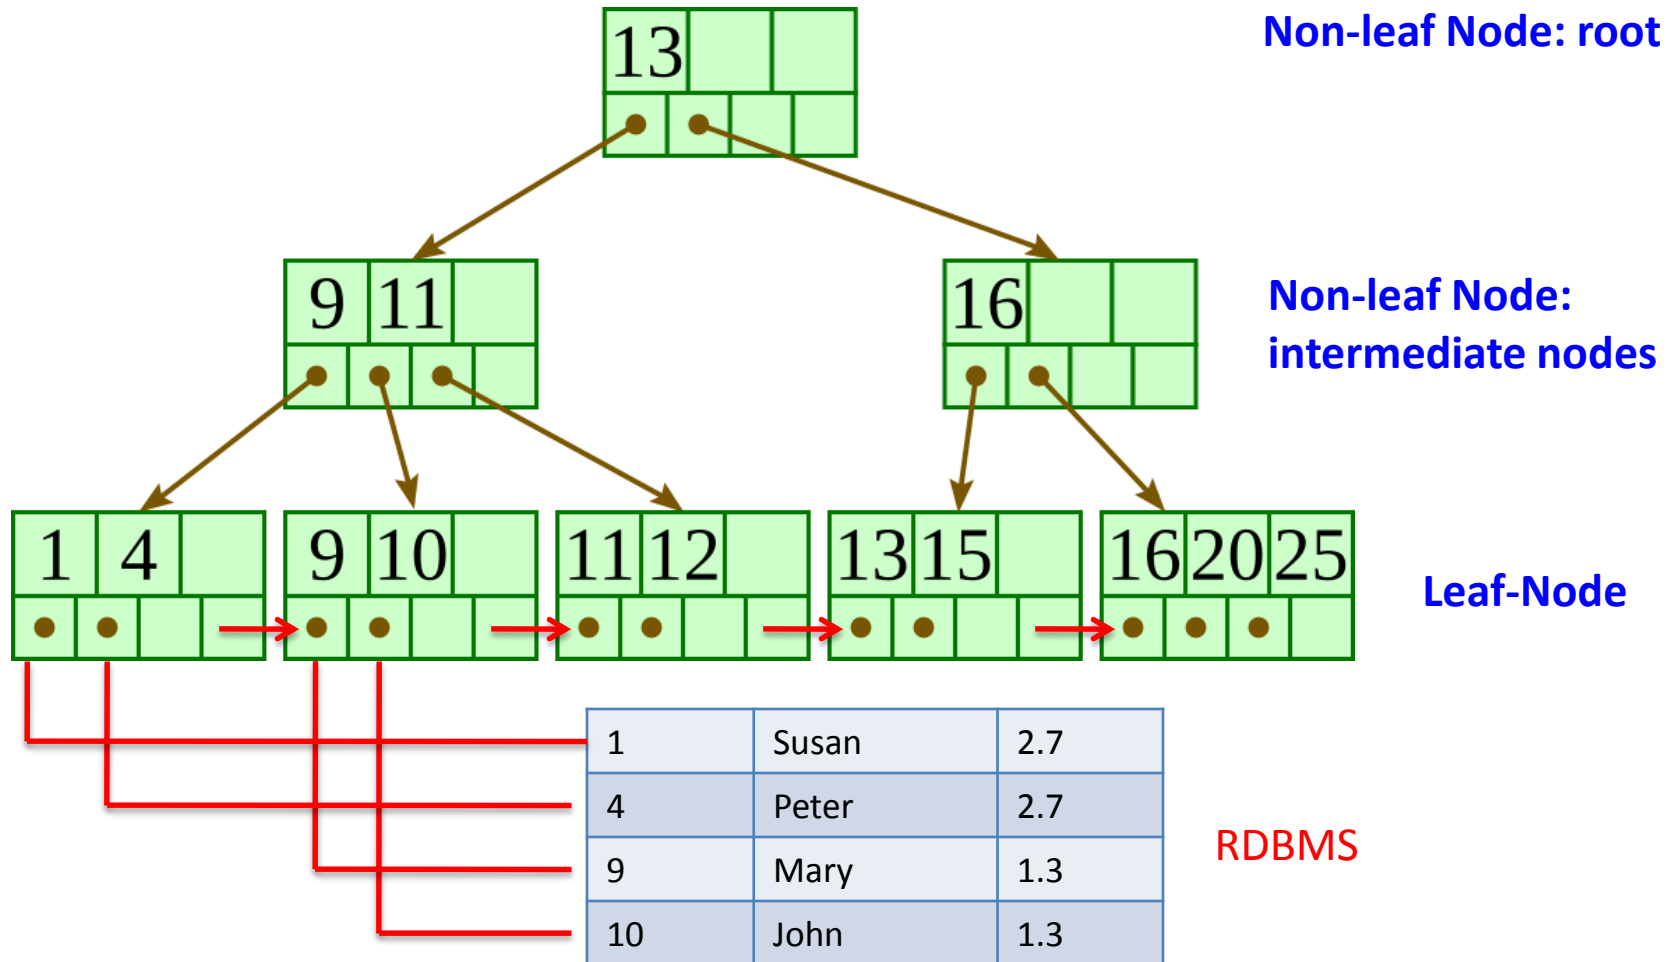
- **B**alanced **I**terative **R**educing and **C**lustering
Using **H**ierarchies
- Why BIRCH?
 - Overcome the bottleneck of datasets not being able to fit in main memory
- Organize the clustering features in CF tree, which is structurally similar to B+ tree

B+ Tree

- A popular index structure in Relational Database Management System
- Advantage
 - Suitable for dynamic updates
 - Balanced
 - Minimum space usage guarantee (50%)


B+ Tree

$n = 4$



B+ Tree

$n = 4$



Non-leaf Node: root

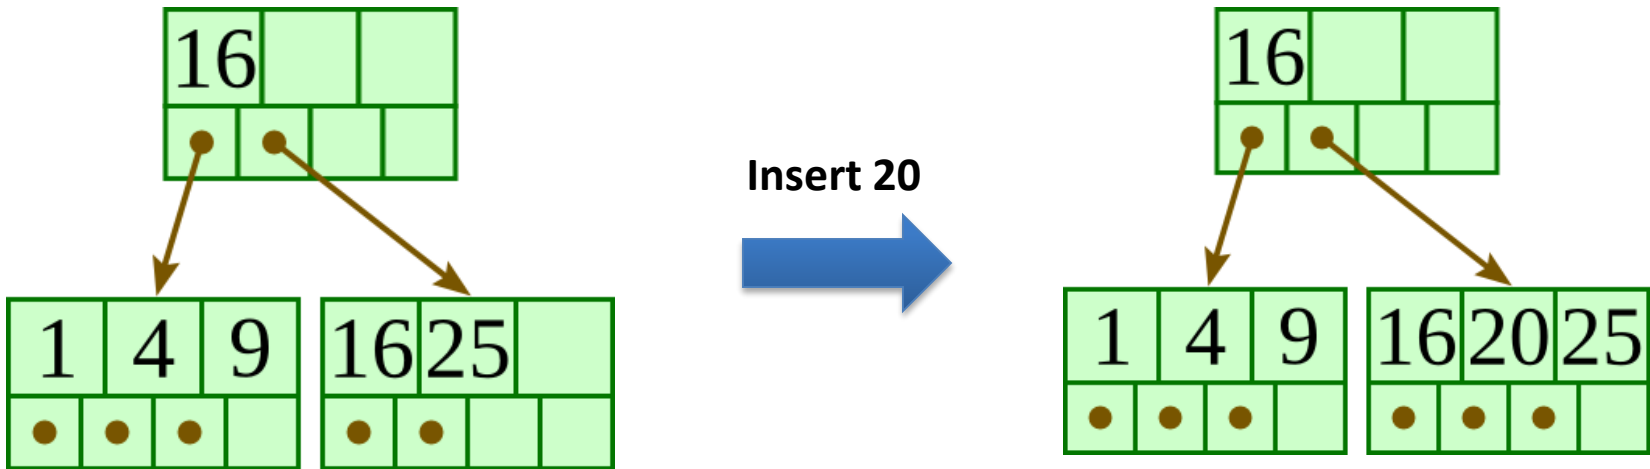
Non-leaf Node:
intermediate nodes

Leaf-Node

- B+ Tree is Balanced
 - The depth is the same for each path from the root to a leaf
 - Every node except the root must be at least half full
 - In this case, each intermediate node and leaf node must have 2 pointers


B+ Tree: Insertion

- Case 1: the node has empty space




B+ Tree: Insertion



- Case 2: the node is full



Insert 13



Insert 12



B+ Tree

- More about B+ Tree

- Visualization

Tool <https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

BIRCH

- Clustering Features: A summary of statistics of the cluster
 - CF = (N, LS, SS)
 - N – Number of data points
 - LS – Linear sum of N points $\sum_{i=1}^N X_i$
 - SS – Square sum of N points $\sum_{i=1}^N X_i^2$
 - A CF entry has sufficient information to calculate the centroid, radius, diameter, and other distance measures

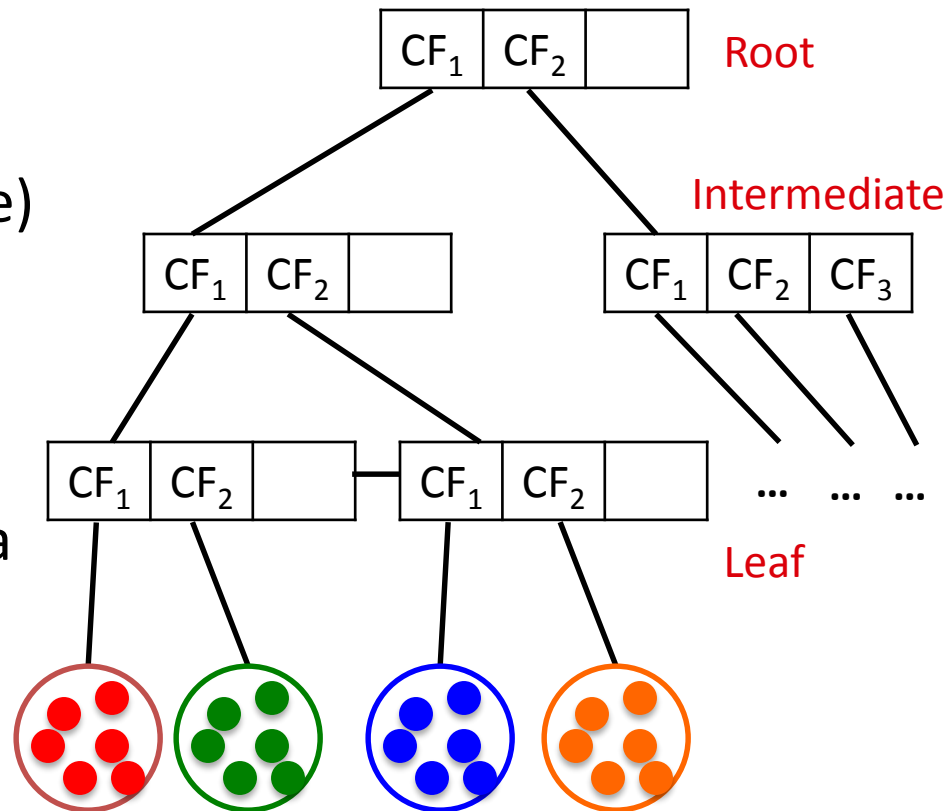
BIRCH

- CF Tree

- Parameters

- **B** = branching factor (max children in a non-leaf node)
 - **L** = number of entries in leaf node
 - **T** = threshold for diameter or radius of the cluster in a leaf

- CF entry in parent = sum of CF entries of a child of that entry



BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster tightness threshold = T




Leaf node with two entries

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster tightness threshold = T



- Object 3 is closer to entry 1
- However, adding object 3 exceed Cluster 1 threshold

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster Tightness Threshold = T



Leaf node with three entries

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster tightness threshold = T



- Object 4 is closer to entry 3
- Cluster 2 remains compact after adding object 4

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster Tightness Threshold = T



Add to cluster 2; Update CF entry

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster Tightness Threshold = T



- Object 5 is closer to entry 2
- However, adding object 5 exceed Cluster 3 threshold
- Exceeds the limit for number of entries in leaf node

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster tightness threshold = T



Split the leaf node; Propagate CF entries one level up

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster tightness threshold = T




- Object 6 is closer to entry 1.2
- Cluster 3 remains compact

BIRCH

- Building the CF tree


Data Objects



Branching factor = 2

Number of entries in leaf node = 3

Cluster tightness threshold = T




Add to cluster 3, Update CF entry

Density-based clustering

- Features:
 - Discover clusters of arbitrary shape.
 - Handle noise.
 - One scan.


DBSCAN

- Inputs:
 - (1) Radius: 1 cm
 - (2) Minimum # of neighbors: 3
- Identify Three kinds of objects:
 - (1) Core object (red)
 - (2) Outlier (blue)
 - (3) Border object (yellow)



Basic Concepts: ε -neighborhood & core objects


- ▶ The neighborhood within a radius ε of a given object is called the **ε -neighborhood** of the object



- ▶ If the ε -neighborhood of an object contains at least a minimum number, **MinPts**, of objects then the object is called a **core object**


→ **Example:** $\varepsilon = 1 \text{ cm}$, $\text{MinPts}=3$

m and **p** are core objects because their ε -neighborhoods contain at least 3 points



Directly density-Reachable Objects

- ▶ An object **p** is **directly density-reachable** from object **q** if **p** is within the ε -neighborhood of **q** and **q** is a core object




→ **Example:**

q is directly density-reachable from **m**

m is directly density-reachable from **p**
and vice versa

Density-Reachable Objects

- ▶ An object **p** is **density-reachable** from object **q** with respect to ϵ and **MinPts** if there is a chain of objects p_1, \dots, p_n where $p_1 = q$ and $p_n = p$ such that p_{i+1} is directly reachable from p_i with respect to ϵ and MinPts




→ Example:

q is density-reachable from **p** because **q** is directly density-reachable from **m** and **m** is directly density-reachable from **p**

p is not density-reachable from **q** because **q** is not a core object

Density-Connectivity

- ▶ An object **p** is **density-connected** to object **q** with respect to ϵ and **MinPts** if there is an object **o** such as both **p** and **q** are density reachable from **o** with respect to ϵ and MinPts



→ **Example:**

p, q and **m** are all density connected


DBSCAN algorithm

- (1) Arbitrary select a point p .
- (2) Retrieve all points density-reachable from p .
- (3) If p is a core point, a cluster is formed.
- (4) If p is a border point, no points are density-reachable from p and DBSCAN visits the next point of the database.
- (5) Continue the process until all the points have been processed.


DBSCAN algorithm

1


MinPts=3




2



3




4




DBSCAN algorithm


5



6



7



8

