



Sequence Clustering

CS 145
Fall 2015

ApproxMAP

- Sequential Pattern Mining
- Support Framework
- Multiple Alignment Framework
- Evaluation
- Conclusion

Inherent Problems

- **Exact match**
 - A pattern gets support from a sequence in the database if and only if **the pattern is exactly contained in the sequence**
 - Often may **not find general long patterns** in the database
 - For example, many customers may share similar buying habits, but few of them follow an exactly same pattern
- **Mines complete set: Too many trivial patterns**
 - Given long sequences with noise
 - ❖ **too expensive and too many patterns**
 - Finding max / closed sequential patterns is non-trivial
 - ❖ **In noisy environment, still too many max/close patterns**

⇒ **Not Summarizing Trend**

Multiple Alignment

- line up the sequences to detect the trend
 - Find common patterns among strings
 - DNA / bio sequences

P	A	T	T	T	E	R	N
P	A	O	O	T	E	R	M
P	O	O	T	T	O	R	N
O	A	O	T	T	E	R	B
P	O	S	Y	Y	R	T	N
P	A	O	T	T	E	R	N

Edit Distance

- ✓ Pairwise Score = edit distance = $\text{dist}(S_1, S_2)$
 - Minimum # of ops required to change S_1 to S_2
 - Ops = INDEL(a) and/or REPLACE(a,b)

P	A	T	T	T	E	R	N
P	A	0	0	T	E	R	M
		INDEL	INDEL				REPL

- **Multiple Alignment Score**
 - $\sum \text{PS}(\text{seq}_i, \text{seq}_j) \quad (\forall 1 \leq i \leq N \text{ and } 1 \leq j \leq N)$
 - **Optimal alignment : minimum score**

Weighted Sequence

- **Weighted Sequence : profile**
 - **Compress a set of aligned sequences into one sequence**

seq_1	(A)		(B)	(DE)	
seq_2	(AE)	(H)	(BC)	(E)	
seq_3	(A)		(BCG)	(D)	
Weighted Sequence	(A:3,E:1):3	(H:1): 1	(B:3,C:2, G:1):3	(D:2, E:2):3	3

Consensus Sequence

- $\text{strength}(i, j) = \frac{\text{\# of occurrences of item } i \text{ in position } j}{\text{total \# of sequences}}$
- Consensus itemset (j)
 - $\{i_a \mid \forall i_a \in (I \cup ()) \ \& \ \text{strength}(i_a, j) \geq \text{min_strength}\}$
- Consensus sequence : **min_strength=2**
 - concatenation of the consensus itemsets for all positions excluding any null consensus itemsets

seq_1	(A)		(B)	(DE)	
seq_2	(AE)	(H)	(BC)	(E)	
seq_3	(A)		(BCG)	(D)	
Weighted Sequence	(A:3,E:1):3	(H:1): 1	(B:3,C:2, G:1):3	(D:2, E:2):3	3
Consensus Sequence	(A)		(BC)	(DE)	

Multiple Alignment Pattern Mining

- **Given**
 - N sequences of sets,
 - Op costs (INDEL & REPLACE) for itemsets, and
 - Strength threshold for consensus sequences
 - ❖ can specify different levels for each partition
- **To**
 - (1) **partition** the N sequences into K sets of sequences such that the sum of the K multiple alignment scores is minimum, and
 - (2) find the optimal **multiple alignment** for each partition, and
 - (3) find the **pattern consensus sequence and the variation consensus sequence for each partition**

ApproxMAP

(Approximate Multiple Alignment Pattern mining)

- **Exact solution : Too expensive!**
- **Approximation Method**
 - **Group : $O(kN) + O(N^2L^2I)$**
 - ❖ partition by Clustering (k-NN)
 - ❖ distance metric
 - **Compress : $O(nL^2)$**
 - ❖ multiple alignment (greedy)
 - **Summarize : $O(1)$**
 - ❖ Pattern and Variation Consensus Sequence
 - **Time Complexity : $O(N^2L^2I)$**

Multiple Alignment : Weighted Sequence

seq₃	(A)		(B)	(DE)		
seq₂	(AE)	(H)	(B)	(D)		
WS₁	(A:2,E:1):2	(H:1):1	(B:2):2	(D:2,E:1):2		2
seq₄	(A)		(BCG)	(D)		
WS₂	(A:3,E:1):3	(H:1):1	(B:3,C:1,G:1):3	(D:3,E:1):3		3

Evaluation Method: Criteria & Datasets

- **Criteria**

- **Recoverability : max patterns**
 - ❖ degree of the underlying patterns in DB detected
 - ❖ $\sum E(F_B) * [\max_{\text{res pat } B} (|B \otimes P|) / E(L_B)]$
 - ❖ Cutoff so that $0 \leq R \leq 1$
- # of spurious patterns
- # of redundant patterns
- **Degree of extraneous items in the patterns**
 - ❖ total # of extraneous items in P / total # of items in P



- **Datasets**

- **Random data : Independence between and across itemsets**
- **Patterned data : IBM synthetic data (Agrawal and Srikant)**
- **Robustness w.r.t. noise : alpha (Yang – SIGMOD 2002)**
- **Robustness w.r.t. random sequences (outliers)**



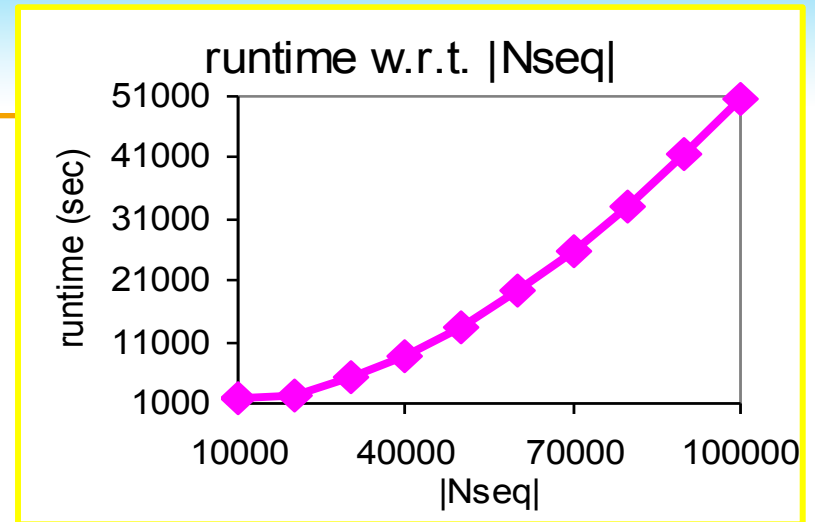

Evaluation : Comparison

	ApproxMAP	Support Framework
Random Data	No patterns with more than 1 item returned	Lots of spurious patterns
Patterned Data 10 patterns embedded into 1000 seqs	<u>k=6 & MinStrgh=30%</u> Recoverability : 92.5% 10 patterns returned 2 redundant patterns 0 spurious patterns 0 extraneous items	<u>MinSup=5%</u> Recoverability : 91.6% 253,924 patterns returned 247,266 redundant patterns 6,648 spurious patterns 93,043=5.2% extraneous items
Noise	Robust	Not Robust Recoverability degrades fast
Outliers	Robust	Robust

Robustness w.r.t. noise



Results : Scalability



Evaluation : Real data

- Successfully applied ApproxMAP to sequence of monthly social welfare services given to clients in North Carolina
- Found interpretable and useful patterns that revealed information from the data

Conclusion : why does it work well?

- **Robust on random & weak patterned noise**
 - Noises can almost never be aligned to generate patterns, so they are ignored
 - If some alignment is possible, the pattern is detected
- **Very good at organizing sequences**
 - when there are “enough” sequences with a certain pattern, they are clustered & aligned
 - When aligning, we start with the sequences with the least noise and add on those with progressively more noise
 - This builds a center of mass to which those sequences with lots of noise can attach to
- **Long sequence data that are not random have unique signatures**

Conclusion

- **Works very well with market basket data**
 - High dimensional
 - Sparse
 - Massive outliers
- **Scales reasonably well**
 - Scales very well w.r.t # of patterns
 - k : scales very well = $O(1)$
 - DB : scales reasonably well = $O(N^2 L^2 I)$

CLUSEQ

- **The primary structures of many biological (macro)molecules are “letter” sequences despite their 3D structures.**
 - **Protein has 20 amino acids.**
 - **DNA has an alphabet of four bases {A, T, G, C}**
 - **RNA has an alphabet {A, U, G, C}**
- **Text document**
- **Transaction logs**
- **Signal streams**
- **Structural similarities at the sequence level often suggest a high likelihood of being functionally/semantically related.**

Problem Statement

- **Clustering based on structural characteristics can serve as a powerful tool to discriminate sequences belonging to different functional categories.**
 - The goal is to create a grouping of sequences such that sequences in each group have similar features.
 - The result can potentially reveal unknown structural and functional categories that may lead to a better understanding of the nature.
- **Challenge: how to measure the structural similarity?**

Measure of Similarity

- **Edit distance:**
 - computationally inefficient
 - only captures the optimal global alignment but ignore many other local alignments that often represent important features shared by the pair of sequences.
- **q -gram based approach:**
 - ignores sequential relationship (e.g., ordering, correlation, dependency, etc.) among q -grams
- **Hidden Markov model:**
 - capture some low order correlations and statistics
 - vulnerable to noise and erroneous parameter setting
 - computationally inefficient

Measure of Similarity



- **Probabilistic Suffix Tree**
 - Effective in capturing significant structural features
 - Easy to compute and incrementally maintain
- **Sparse Markov Transducer**
 - Allows wild cards

Model of CLUSEQ

- **CLUSEQ: exploring significant patterns of sequence formation.**
 - Sequences belonging to one group/cluster may subsume to the same probability distribution of symbols (conditioning on the preceding segment of a certain length), while different groups/clusters may follow different underlying probability distributions.
 - By extracting and maintaining significant patterns characterizing (potential) sequence clusters, one can easily determine whether a sequence should belong to a cluster by **calculating the likelihood of (re)producing the sequence under the probability distribution that characterizes the cluster.**

Model of CLUSEQ

Sequence: $\sigma = s_1 s_2 \dots s_l$

Cluster S :
$$P_S(\sigma) = P_S(s_1) \times P_S(s_2 | s_1) \times \dots \times P_S(s_l | s_1 \dots s_{l-1})$$
$$= \prod_{i=1}^l P_S(s_i | s_1 \dots s_{i-1})$$

Random process:
$$P^r(\sigma) = P^r(s_1) \times P^r(s_2) \times \dots \times P^r(s_l)$$
$$= \prod_{i=1}^l P^r(s_i)$$

If $P_S(\sigma) \gg P^r(\sigma)$, we may consider σ a member of S

Model of CLUSEQ

- **Similarity between σ and S**

$$sim_S(\sigma) = \frac{P_S(\sigma)}{P^r(\sigma)} = \frac{\prod_{i=1}^l P_S(s_i | s_1 \dots s_{i-1})}{\prod_{i=1}^l p(s_i)} = \prod_{i=1}^l \left(\frac{P_S(s_i | s_1 \dots s_{i-1})}{p(s_i)} \right)$$

- **Noise may be present.**
- **Different portions of a (long) sequence may subsume to different conditional probability distributions.**

$$SIM_S(\sigma) = \max_{1 \leq i \leq j \leq l} sim_S(s_i \dots s_j)$$

Model of CLUSEQ

- Give a sequence $\sigma = s_1 s_2 \dots s_l$ and a cluster S , a dynamic programming method can be used to calculate the similarity $SIM_S(\sigma)$. Via a single scan of σ . Let

$$X_i = \frac{P_S(s_i | s_1 \dots s_{i-1})}{p(s_i)}$$

$$Y_i = \max_{1 \leq j \leq i} sim_S(s_j \dots s_i)$$

$$Z_i = \max_{1 \leq i_1 \leq i_2 \leq i} sim_S(s_{i_1} \dots s_{i_2})$$

- Intuitively, X_i , Y_i , and Z_i can be viewed as the similarity contributed by the symbol on the i th position of σ (i.e., s_i), the maximum similarity possessed by any segment ending at the i th position, and the maximum similarity possessed by any segment ending prior to or on the i th position, respectively.

Model of CLUSEQ

- Then, $SIM_S(\sigma) = Z_p$, which can be obtained by

$$Y_i = \max\{Y_{i-1} \times X_i, X_i\}$$

$$Z_i = \max\{Z_{i-1}, Y_i\}$$

$$Y_1 = Z_1 = X_1$$

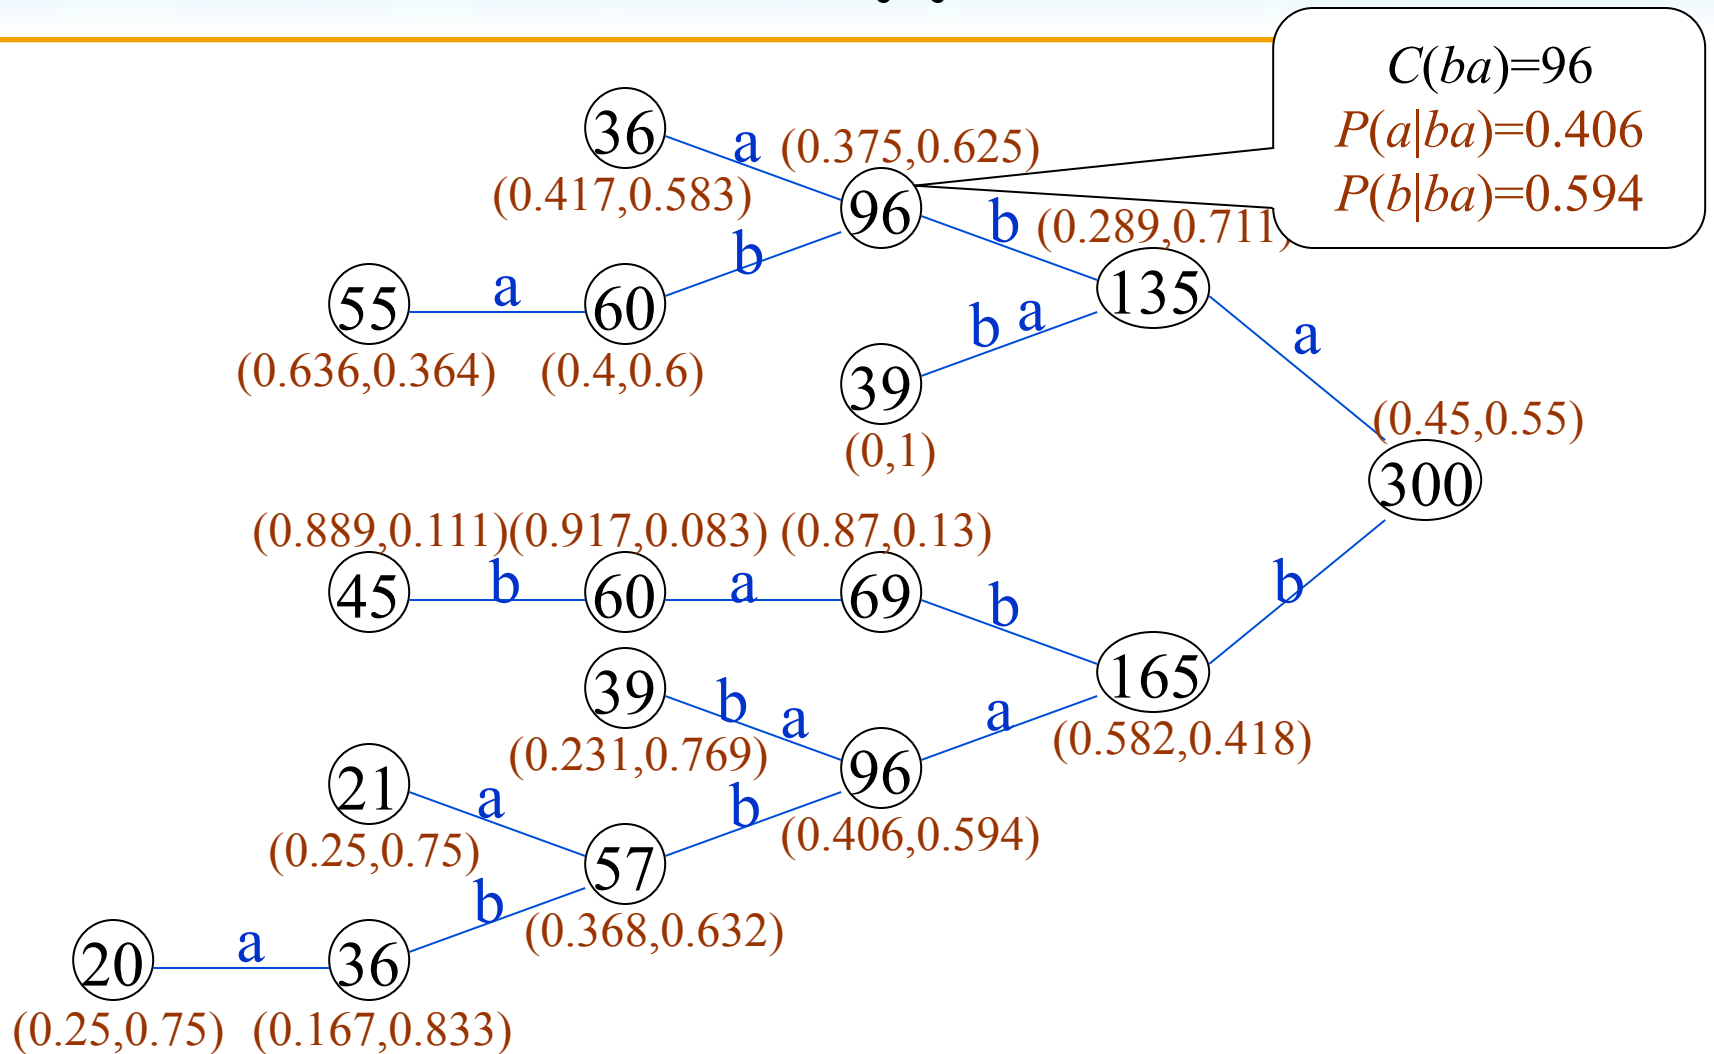
- For example, $SIM_S(bbaa) = 2.10$ if $p(a) = 0.6$ and $p(b) = 0.4$.

Sequence	<i>b</i>	<i>b</i>	<i>a</i>	<i>a</i>
$P_S(s_i s_1 \dots s_{i-1})$	0.55	0.418	0.87	0.406
X_i	1.38	1.05	1.45	0.677
Y_i	1.38	1.45	2.10	1.42
Z_i	1.38	1.45	2.10	2.10

Probabilistic Suffix Tree

- a compact representation to organize the derived CPD for a cluster
- built on the *reversed* sequences
- Each node corresponds to a segment, σ , and is associated with a counter $C(\sigma)$ and a probability vector $P(s_i | \sigma)$.

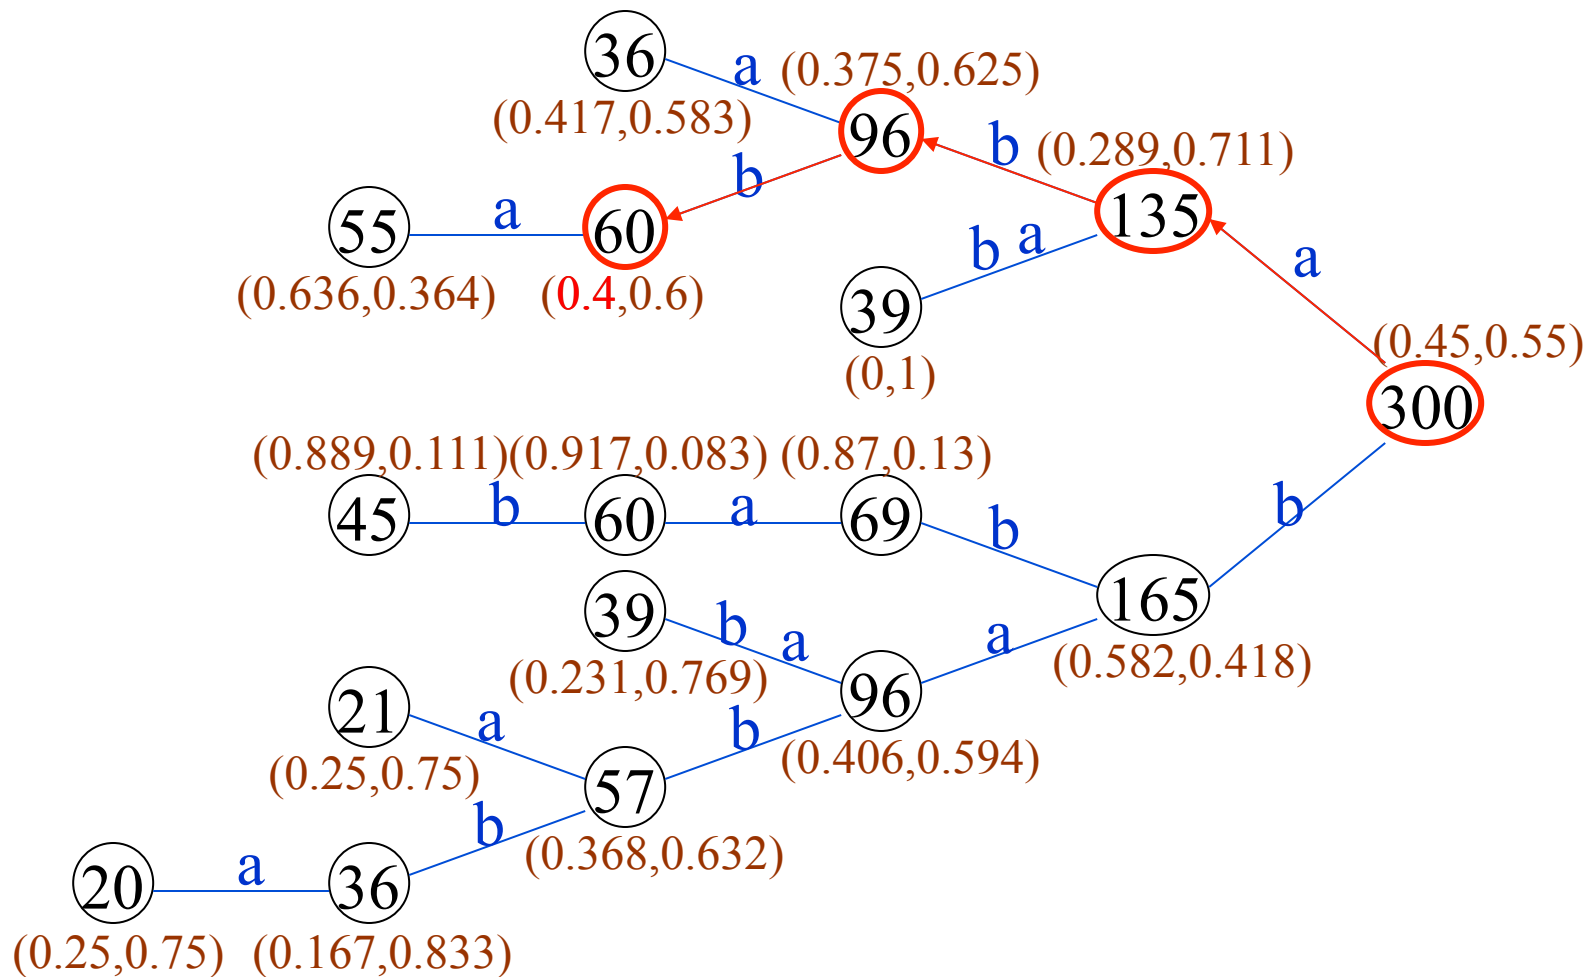
Probabilistic Suffix Tree



Model of CLUSEQ

- Retrieval of a CPD entry $P(s_i | s_1 \dots s_{i-1})$
- The *longest suffix* $s_j \dots s_{i-1}$
 - can be located by traversing from the root along the path “ $\rightarrow s_{i-1} \rightarrow \dots \rightarrow s_2 \rightarrow s_1$ ” until we reach either the node labeled with $s_1 \dots s_i$ or a node where no further advance can be made.
 - takes $O(\min\{i, h\})$ where h is the height of the tree.
- Example: $P(a | bbba)$

$$P(a|bbba) \approx P(a|bba) = 0.4$$



CLUSEQ

- **Sequence Cluster:** a set of sequences S is a sequence cluster if, for each sequence σ in S , the similarity $SIM_S(\sigma)$ between σ and S is greater than or equal to some similarity threshold t .
- **Objective:** automatically group a set of sequences into a set of *possibly overlapping* clusters.


Algorithm of CLUSEQ

Unclustered sequences

- **An iterative process**

- Each cluster is represented by a probabilistic suffix tree.
- The optimal number of clusters and the number of outliers allowed can be adapted by CLUSEQ automatically

- ❖ new cluster generation, cluster split, and cluster consolidation
- ❖ adjustment of similarity threshold



New Cluster Generation

- New clusters are generated from *un-clustered* sequences at the beginning of **each** iteration.

➤ $k' \times f$ new clusters


number of consolidated clusters

$$f = \frac{\max \{k'_n - k'_c, 0\}}{k'_n}$$

k'_n


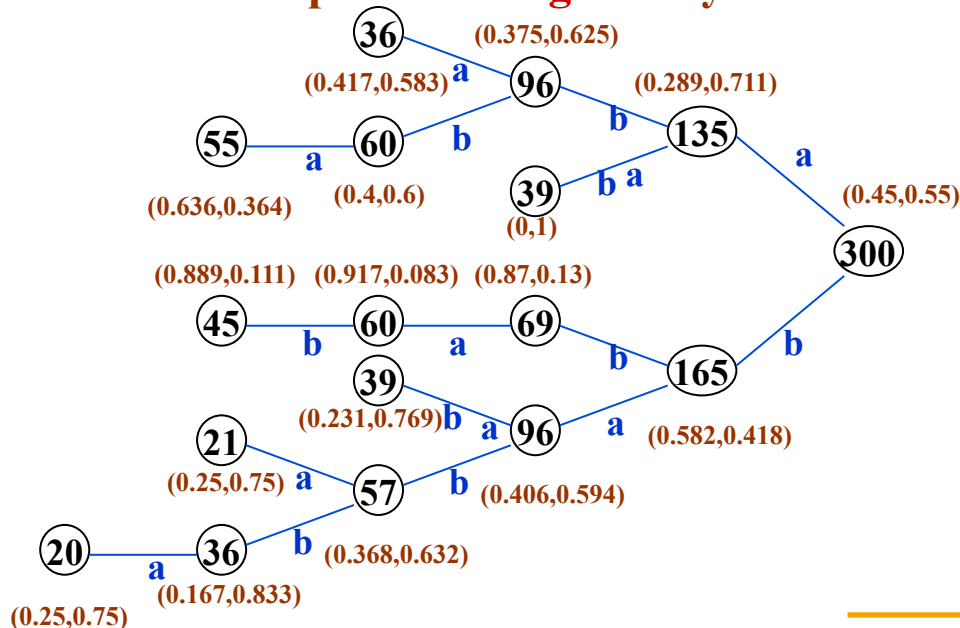
number of clusters

number of new clusters generated at the previous iteration




Sequence Re-Clustering

- For each (sequence, cluster) pair
 - Calculate similarity
 - PST update if necessary
 - ❖ Only **similar portion** is used
 - ❖ The update is **weighted** by the similarity value



Cluster Split

- Check the convergence of each existing cluster
 - Imprecise probabilities are used for each probability entry in PST
 - Split non-convergent cluster



Imprecise Probabilities

- Imprecise probabilities uses two values (p_1, p_2) (instead of one) for a probability.
 - p_1 is called lower probability and p_2 is called upper probability.
 - The true probability lies somewhere between p_1 and p_2 .
 - $p_2 - p_1$ is called imprecision.

Update Imprecise Probabilities

- Assuming the prior knowledge of a (conditional) probability is (p_1, p_2) and the occurrences in the new experiment is a out of b trials.

$$p'_1 = \frac{a + s \times p_1}{b + s} \qquad p'_2 = \frac{a + s \times p_2}{b + s}$$



where s is the learning parameter which controls the weight that each experiment carries.

Properties

- The following two properties are very important.
 - If the probability distribution stays static, then p_1 and p_2 will converge to the true probability.
 - If the experiment agrees with the prior assumption, the range of imprecision decreases after applying the new evidence, e.g., $p_2' - p_1' < p_2 - p_1$.
- The clustering process terminates when the imprecision of all significant nodes is less than a small threshold.



Cluster Consolidation

- Starting from the smallest cluster
- Dismiss clusters that have few sequence **not covered** by other clusters



Adjustment of Similarity Threshold

- Find the sharpest turn of the similarity distribution function



Algorithm of CLUSEQ

- **Implementation issues**
 - **Limited memory space**
 - ❖ Prune the node with smallest count first.
 - ❖ Prune the node with longest label first.
 - ❖ Prune the node with expected probability vector first.
 - **Probability smoothing**
 - ❖ Eliminates zero empirical probability
 - **Other considerations**
 - ❖ Background probabilities
 - ❖ A priori knowledge
 - ❖ Other structural features

Experimental Study

- We have experimented with a protein database of 8000 proteins from 30 families from SWISS-PROT database.

Model	CLUSEQ	Edit Distance	Edit Distance with Block Operations	Hidden Markov Model	Q-gram
Accuracy	92%	23%	90%	91%	75%
Response Time (sec)	144	487	13754	3117	132

Experimental Study

Synthetic data

Initial t	1.05	1.5	2	3
Final t	1.99	2.01	2	1.99
Response time	8011	7556	6754	7234
precision	81.3%	83.1%	83.4%	81.9%
recall	82.1%	82.8%	83.6%	82.7%

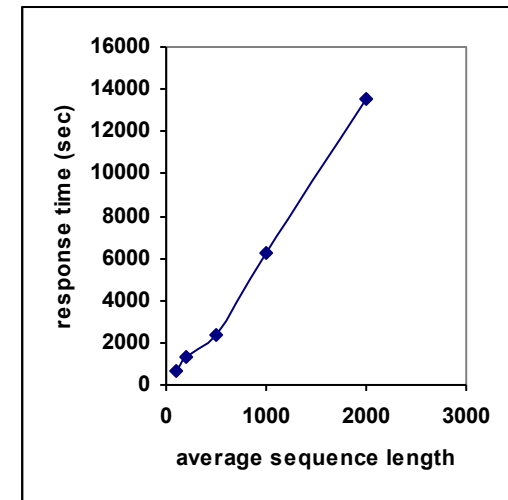
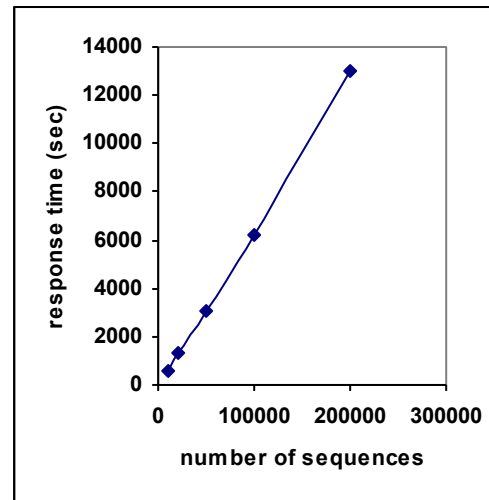

Experimental Study

Synthetic data

Initial cluster number	1	20	100	200
Final cluster number	102	99	101	102
Response time	10112	9023	6754	8976
precision	81.3%	82.1%	82.6%	81%
recall	81.6%	82%	83.4%	81.7%

Experimental Study

- **CLUSEQ has linear scalability with respect to the number of clusters, number of sequences, and sequence length.**



Remarks

- **Similarity measure**
 - Powerful in capturing high order statistics and dependencies
 - Efficient in computation — linear complexity
 - Robust to noise
- **Clustering algorithm**
 - High accuracy
 - High adaptability
 - High scalability
 - High reliability

References

- CLUSEQ: efficient and effective sequence clustering, *Proceedings of the 19th IEEE International Conference on Data Engineering (ICDE)*, 2003.
- A frame work towards efficient and effective protein clustering, *Proceedings of the 1st IEEE CSB*, 2002.